

Serial Protocol Decode — user manual

Ian Ameline · version 1.05 — **beta** · MIT licence

This app turns a Keithley DMM6500 into a serial decoder. Clip onto a UART line, press **Capture**, and read the bytes on the front panel. You do not have to tell it the baud rate, the frame format or which way up the signal is — it works all of that out from the signal itself.

Three things make it worth reaching for over a scope's decode: the **window is large** (a screenful is ~240 bytes, and a recording holds 32 768), every capture is **appended to the USB key automatically** while a key is in the slot, so you have a log without asking for one, and the **serial parameters are detected rather than configured**. For most debugging that is the whole job — connect, press, read.

This is a beta. Every number in this manual was measured on a real DMM6500 against a real signal generator unless the text says otherwise — and where something is calculated rather than measured, or untested, it says so in that spot. What that cannot cover is the variety of real devices, so if the app tells you something you can prove is wrong, that is the bug worth reporting. It is built so that **refusing with a reason is acceptable and confident nonsense is not**, and anything it is unsure about it says on the note row at the bottom of the screen.

Hooking it up — read this part

The line must stay between -10 V and +10 V. The app fixes the meter on its 10 V range. 3.3 V CMOS and 5 V TTL are both fine as they are. For a 12 V line, put a 2:1 resistor divider in front of it — and if it is an open-drain bus, pick resistors big enough not to load it down.

Connect INPUT LO to the ground of the circuit you are probing. Everything is measured at INPUT HI relative to INPUT LO, so without a shared ground the readings mean nothing.

Do not put a differential bus straight across HI and LO. RS-485, CAN and LIN need to go through a transceiver first, or you need to tap one side of the pair against ground.

A DC offset does not matter. The app measures the two voltage levels on your line and puts its decision threshold between them, so it does not care where they sit — only that both fit inside ± 10 V. A 1.6 V swing sitting up at 6 V decodes fine. **1.6 V is the smallest swing tested.** A 60 mV swing is refused outright.

Quick start

1. Connect the line to **INPUT HI** and its ground to **INPUT LO**.
2. Press **Capture** while the device is talking.
3. Read the bytes. Press **View** to switch between hex and plain text.

That is the whole normal workflow. The first capture also locks the baud rate when it is confident, so later captures are faster.

Two more things worth knowing before you need them: **Mode** turns Capture into a recording of up to 32 768 bytes straight to the USB key, still one press — and once a long job starts, **nothing stops it early**, so decide the size before you press — see **The TRIGGER key** below.

The screen

SERIAL DECODE - 240B FRAME

End App

BAUD	FORMAT	IDLE	LOGIC	THRESH	SA/BIT	RATE	BYTES	ERR	FIT
19200	7E1	HIGH	3V3 CMOS	1.63 V	8.3	160 kS/s	235	0	1.00

Start bit

0000

20 64 6F 67 2E 20 54 48

45 20 51 55 49 43 4B 20

0016

42 52 4F 57 4E 20 46 4F

58 20 4A 55 4D 50 53 20

0032

4F 56 45 52 20 54 48 45

20 4C 41 5A 59 20 44 4F

0048

47 2E 20 30 31 32 33 34

35 36 37 38 39 20 21 22

0064

23 24 25 26 27 28 29 2A

2B 2C 2D 2E 2F 3A 3B 3C

0080

3D 3E 3F 40 5B 5C 5D 5E

5F 60 7B 7C 7D 7E 74 68

0096

65 20 71 75 69 63 6B 20

62 72 6F 77 6E 20 66 6F

0112

78 20 6A 75 6D 70 73 20

6F 76 65 72 20 74 68 65

0128

20 6C 61 7A 79 20 64 6F

67 2E 20 54 48 45 20 51

0144

55 49 43 4B 20 42 52 4F

57 4E 20 46 4F 58 20 4A

0160

55 4D 50 53 20 4F 56 45

52 20 54 48 45 20 4C 41

0176

5A 59 20 44 4F 47 2E 20

30 31 32 33 34 35 36 37

0192

38 39 20 21 22 23 24 25

26 27 28 29 2A 2B 2C 2D

0208

2E 2F 3A 3B 3C 3D 3E 3F

40 5B 5C 5D 5E 5F 60 7B

0224

7C 7D 7E 74 68 65 20 71

75 69 63

dog. THE QUICK
BROWN FOX JUMPS
OVER THE LAZY DO
G. 0123456789 !"#\$%&'()*+,-./:;<
=>?@[\\]^_`{|}~th
e quick brown fo
x jumps over the
lazy dog. THE Q
UICK BROWN FOX J
UMPS OVER THE LA
ZY DOG. 01234567
89 !"#\$%&'()*+,-
./:;<=>?@[\\]^_`{|
}|~the quic

FRAME HEX pg 1/1 bytes 1-235/235 win 240 [done]log: bytes324 472 B

Capture

View

Mode

NewLog

Save

Options

Along the top are the things the app worked out about your line:

Field	Meaning
BAUD	The rate it found. A padlock means it is locked and will not be re-detected.
FORMAT	Data bits, parity, stop bits — 8N1 , 7E1 and so on.
IDLE	What the line rests at: HIGH for CMOS/TTL, LOW for an RS-232-style line.
LOGIC	The family it looks like, e.g. 3V3 CMOS , 5V TTL .
THRESH	The voltage it is using to tell 0 from 1.
SA/BIT	Samples per bit. More is better; 8 is comfortable, 4 still works.
RATE	How fast the meter sampled.
BYTES / ERR	How many bytes came out, and how many of them you cannot trust.
FIT	How well one bit-length explains what it measured, from 0 to 1. Below about 0.9 means the timing was messy.

A dash means "not measured yet".

Broken bytes are shown, not hidden. ERR turns red when there are any, and every row containing one is drawn red across its whole width — offset, hex and ASCII together — so damage is found by looking rather than by counting. Flagged bytes read as ? in the ASCII gutter:

SERIAL DECODE - 240B FRAME

End App

BAUD	FORMAT	IDLE	LOGIC	THRESH	SA/BIT	RATE	BYTES	ERR	FIT
76800	7E1	HIGH	3V3 CMOS	1.63 V	8.3	635 kS/s	236	15	1.00

Start bit

began mid-byte -- the first 15 bytes are misaligned.

0000	12 29 4D 6D 09 2D 49 69	4D 5A 63 6A 73 7B 02 5B	??M???I?M???s??[
0016	5C 5D 5E 5F 60 7B 7C 7D	7E 74 68 65 20 71 75 69	\]^_`{ }~the qui
0032	63 6B 20 62 72 6F 77 6E	20 66 6F 78 20 6A 75 6D	ck brown fox jum
0048	70 73 20 6F 76 65 72 20	74 68 65 20 6C 61 7A 79	ps over the lazy
0064	20 64 6F 67 2E 20 54 48	45 20 51 55 49 43 4B 20	dog. THE QUICK
0080	42 52 4F 57 4E 20 46 4F	58 20 4A 55 4D 50 53 20	BROWN FOX JUMPS
0096	4F 56 45 52 20 54 48 45	20 4C 41 5A 59 20 44 4F	OVER THE LAZY DO
0112	47 2E 20 30 31 32 33 34	35 36 37 38 39 20 21 22	G. 0123456789 !"#\$%
0128	23 24 25 26 27 28 29 2A	2B 2C 2D 2E 2F 3A 3B 3C	%&'()*+,-./:;<=>?
0144	3D 3E 3F 40 5B 5C 5D 5E	5F 60 7B 7C 7D 7E 74 68	@[\]^_`{ }~the
0160	65 20 71 75 69 63 6B 20	62 72 6F 77 6E 20 66 6F	e quick brown fo
0176	78 20 6A 75 6D 70 73 20	6F 76 65 72 20 74 68 65	x jumps over the
0192	20 6C 61 7A 79 20 64 6F	67 2E 20 54 48 45 20 51	lazy dog. THE Q
0208	55 49 43 4B 20 42 52 4F	57 4E 20 46 4F 58 20 4A	UICK BROWN FOX J
0224	55 4D 50 53 20 4F 56 45	52 20 54 48	UMPS OVER TH

FRAME HEX pg 1/1 bytes 1-236/236 win 242 [done]

log: bytes337 947 B

Capture

View

Mode

NewLog

Save

Options

That is a real 240B capture that began part way through a byte. **The note row names the cause** — began mid-byte -- the first N bytes are misaligned. — which is the difference between a puzzle and a known condition: the app arms on a busy line, so it lands wherever the traffic happens to be, and the bytes before the first gap between messages cannot be framed. On a busy line roughly one capture in eight starts this way.

The note and ERR agree, and that is deliberate. The note's number is how far the damage reaches: the last byte in the opening region the decoder could not frame, so "the first N" is literally true and nothing past N is being accused. ERR is then never less than N. Above, the note says 15 and ERR says 15 — while only 11 of those bytes actually failed a check. The other four are the reason ERR counts the whole region: inside a misaligned head the bit boundaries are in the wrong places, so a frame there can satisfy its parity and stop bit by luck and still hand back a plausible wrong byte. 3, 7, 9 and 13 came out as M, I, M and s — perfectly ordinary characters, and wrong. **ERR counts bytes you cannot trust, not checks that failed**, because the silently wrong ones are the dangerous ones.

Two exclusions remain, and both keep ERR at 0 on good data. When no misaligned head is identified, the first three frames are ignored — a gapless line has no idle for the framer to anchor on, so it routinely resynchronises a frame or two in, and counting that made ERR read 1 or 2 on every healthy capture. The final frame is ignored always, because the capture boundary cuts it in half. Neither exclusion applies across a head the app has evidence for, which is what stopped the panel reading ERR 1 beside a note naming four bytes.

Everything from the first gap onward decoded cleanly, and the log on the USB key holds every byte: one that failed its parity or stop bit is still written, marked, rather than dropped.

Underneath: the **trigger cell**, the bytes, then two rows of text. The **status row** tells you where you are (FRAME HEX pg 1/1 bytes 1-155/155 win 240 [done]) and names the log file.

Which mode you are in is in the title bar — SERIAL DECODE - 240B FRAME , - 8K CAPTURE or - 32K CAPTURE , in the largest type on the screen.

The cell at the left of the note row is what a capture **WAITS for**, colour-coded so a glance is enough: green **Start bit** (the normal case — it waits for the line's own start bit), yellow **Free run** (it does not wait for anything), cyan **Trigger key** (it waits for the physical key).

The cell at the bottom right is the rear **BNC**, and it is **empty when the rear BNC is off** — which is the default. When it is not, it names the direction: green **EXT TRIG IN**, amber **EXT TRIG OUT**, cyan **EXT TRIG I/O**. A trailing **?** means the input is switched on but *not in force*, which happens under **Free run** — free run means do not wait, so there is nothing for the input to gate.

During a recording the status row carries a live counter — **recording... 40 % of the buffer**, then **decoding... 4096/8192 bytes** — with a **cyan progress bar** beside it, and the cell to its right reads **Capture or Mode = Stop**.

View switches the byte display to plain text, which is what you want for anything human-readable — and a row carrying a flagged byte is red here too:

SERIAL DECODE - 240B FRAME

End App

BAUD	FORMAT	IDLE	LOGIC	THRESH	SA/BIT	RATE	BYTES	ERR	FIT
9600	8N1	HIGH	3V3 CMOS	1.63 V	8.3	80 kS/s	237	5	1.00

Start bit began mid-byte -- the first 5 bytes are misaligned.

...? Excepteur sint occaecat ..cupidatat non proident, sunt in culpa qui officina deserunt mollit ..anim id est laborum. Lorem ipsum dolor sit amet, consectetur adipiscing ..elit, sed do eiusmod tempor incididunt ut labore et dLorem ipsu

FRAME TEXT pg 1/1 bytes 1-237/237 win 239 [done]log: bytes339 21217 B

CaptureViewModeNewLogSaveOptions

The note row is the important one. Warnings appear there — an ambiguous baud rate, a line that disagrees with the rate you locked, a recording that stopped early. If more than one applies it ends with **(+2 more)**, and pressing **Save** writes all of them to a file.

Here it is doing that — a capture that started part way through a byte, which the app detected and said so rather than showing two wrong bytes without comment:

SERIAL DECODE - 240B FRAME

End App

BAUD	FORMAT	IDLE	LOGIC	THRESH	SA/BIT	RATE	BYTES	ERR	FIT
76800	7E1	HIGH	3V3 CMOS	1.63 V	8.3	635 kS/s	236	2	1.00

Start bit
began mid-byte -- the first 2 bytes are misaligned.

0000	07 53 55 4D 50 53 20 4F	56 45 52 20 54 48 45 20	??UMPS OVER THE
0016	4C 41 5A 59 20 44 4F 47	2E 20 30 31 32 33 34 35	LAZY DOG. 012345
0032	36 37 38 39 20 21 22 23	24 25 26 27 28 29 2A 2B	6789 !"#\$%&'()*+,-./:;<=>?@[\\]^_`{ }~the quick brown fox jumps over the lazy dog. THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG. 0123456789 !"#\$%&'()*+,-./:;<=>?@[\\]^_`{ }~the quick brown fox jumps over the
0048	2C 2D 2E 2F 3A 3B 3C 3D	3E 3F 40 5B 5C 5D 5E 5F	
0064	60 7B 7C 7D 7E 74 68 65	20 71 75 69 63 6B 20 62	
0080	72 6F 77 6E 20 66 6F 78	20 6A 75 6D 70 73 20 6F	
0096	76 65 72 20 74 68 65 20	6C 61 7A 79 20 64 6F 67	
0112	2E 20 54 48 45 20 51 55	49 43 4B 20 42 52 4F 57	
0128	4E 20 46 4F 58 20 4A 55	4D 50 53 20 4F 56 45 52	
0144	20 54 48 45 20 4C 41 5A	59 20 44 4F 47 2E 20 30	
0160	31 32 33 34 35 36 37 38	39 20 21 22 23 24 25 26	
0176	27 28 29 2A 2B 2C 2D 2E	2F 3A 3B 3C 3D 3E 3F 40	
0192	5B 5C 5D 5E 5F 60 7B 7C	7D 7E 74 68 65 20 71 75	
0208	69 63 6B 20 62 72 6F 77	6E 20 66 6F 78 20 6A 75	
0224	6D 70 73 20 6F 76 65 72	20 74 68 65	

FRAME HEX pg 1/1 bytes 1-236/236 win 242 [done]
log: bytes337 2843 B

Capture

View

Mode

NewLog

Save

Options

An empty note row is not a promise that everything is perfect. It means the app did not detect anything wrong. A noise spike landing inside a data bit changes that byte with nothing to give it away.

The buttons

Button	Does
Capture	Take a capture and decode it. In a recording mode, does the whole recording.
View	Switch the display: hex + ASCII → plain text → back.
Mode	Choose what Capture does: a screenful, an 8 kB recording, or a 32 kB one.
NewLog	Start a new numbered log file on the USB key.
Save	Write a full report of what is on screen to the USB key.
Options	The settings screen.
Up / Dn	Right-hand edge. Only appear when there is more than one page. Once you press either, page N of M appears at the right of the note row.
Lock Rate	Right-hand edge. Only appears when the rate is not locked.

Every button responds immediately except the ones that take a capture. **Capture takes about 1.7 seconds when the rate is locked, and 4 to 9 seconds when it is not** — because to find an unknown rate it has to capture two or three times. Slower lines take longer, not less. The exact time varies from press to press because it depends on where the traffic happens to be when you press.

A press during a capture is not lost — it is queued and acted on when the capture finishes. So if the panel seems to ignore you, wait rather than pressing again.

Paging through a long capture

SERIAL DECODE - 32K CAPTURE

End App

BAUD	FORMAT	IDLE	LOGIC	THRESH	SA/BIT	RATE	BYTES	ERR	FIT
9600	8N1	HIGH	3V3 CMOS	1.63 V	4.1	39 kS/s	32768	6	1.00
Start bit panel shows bytes 24577-32768 of 32768; all are in bytes340.txt (+1 more) page 4 of 35									
25296	61 6D 2C 20 71 75 69 73	20 0D 0A 6E 6F 73 74 72	am, quis ..nostr						
25312	75 64 20 65 78 65 72 63	69 74 61 74 69 6F 6E 20	ud exercitation						
25328	75 6C 6C 61 6D 63 6F 20	6C 61 62 6F 72 69 73 20	ullamco laboris						
25344	6E 69 73 69 20 75 74 20	61 6C 69 71 75 69 70 20	nisi ut aliquip						
25360	65 78 20 65 61 20 63 6F	6D 6D 6F 64 6F 20 0D 0A	ex ea commodo ..						
25376	63 6F 6E 73 65 71 75 61	74 2E 20 44 75 69 73 20	consequat. Duis						
25392	61 75 74 65 20 69 72 75	72 65 20 64 6F 6C 6F 72	aute irure dolor						
25408	20 69 6E 20 72 65 70 72	65 68 65 6E 64 65 72 69	in reprehenderi						
25424	74 20 69 6E 20 76 6F 6C	75 70 74 61 74 65 20 76	t in voluptate v						
25440	65 6C 69 74 20 0D 0A 65	73 73 65 20 63 69 6C 6C	elit ..esse cill						
25456	75 6D 20 64 6F 6C 6F 72	65 20 65 75 20 66 75 67	um dolore eu fug						
25472	69 61 74 20 6E 75 6C 6C	61 20 70 61 72 69 61 74	iat nulla pariat						
25488	75 72 2E 20 45 78 63 65	70 74 65 75 72 20 73 69	ur. Excepteur si						
25504	6E 74 20 6F 63 63 61 65	63 61 74 20 0D 0A 63 75	nt occaecat ..cu						
25520	70 69 64 61 74 61 74 20	6E 6F 6E 20 70 72 6F 69	pidatat non proi						
32 kB 32768 bytes 6 err 68 win -> bytes340.txt [cap] log: on first capture									

A 32 kB recording is far more than one screen, so **Up** and **Dn** appear down the right-hand edge and **page N of M** appears at the right of the note row — the two numbers in white, the words dimmed, so the count reads at a glance. The note beside it says which bytes are on screen and that the whole run is in the file, which is the thing you cannot otherwise tell from a screenful taken out of the middle:

panel shows bytes 24577-32768 of 32768; all are in bytes340.txt

In the figure that note is abbreviated, ending **(+1 more)**: this run also began mid-byte, and two notes do not fit one row. **Save** writes both out in full.

The offset column agrees with it — the rows start at 25296, not at zero — and the status row carries the run's own summary: 32768 bytes, 6 it cannot vouch for, 68 windows, and the file it went to. **That 6 is the same figure as ERR**, not a second opinion about it.

The page count follows the view. Press **View** on that same capture and it becomes 7 pages instead of 35, because a text row holds 80 characters and a hex row holds 16 bytes — 1200 bytes a page against 240:

SERIAL DECODE - 32K CAPTURE

End App

BAUD	FORMAT	IDLE	LOGIC	THRESH	SA/BIT	RATE	BYTES	ERR	FIT
9600	8N1	HIGH	3V3 CMOS	1.63 V	4.1	39 kS/s	32768	6	1.00

Start bit

panel shows bytes 24577-32768 of 32768; all are in bytes3... (+1 more)

page 2 of 7

Up

Dn

32 kB 32768 bytes 6 err 68 win -> bytes340.txt [cap]

log: on first capture

Capture

View

Mode

NewLog

Save

Options

Nothing about the capture changed; only how much of it fits on a screen. The byte range in the note row is the same, and so is the summary.

The TRIGGER key — the one hardware button that matters

The **TRIGGER** key is the physical key on the front panel, to the **right** of the screen — the lowest of that column of buttons. It is not part of the app's own row of buttons. It was intended to do two things nothing on the glass can do; **one of them works** — timing a capture — and the other, stopping a long job, does not. Both are described below, because knowing which is which is the difference between waiting confidently and pressing a key that does nothing.

1. It does not stop a long job, and nothing else does either

A recording runs to its own end. Once you press **Capture** in `8 kB` or `32 kB`, nothing on the instrument will cut it short — not a button on the glass, and not the TRIGGER key. Decide before you press. The status row says so before you commit:

```
32 kB  ready -- 32768 bytes to a file, no stop once started
```

There is deliberately no time estimate. Two different durations apply and only one is knowable here: the time your *device* takes to send that many bytes, and the meter's decode, which costs several times more — 8 kB is a couple of seconds on the wire and around half a minute to decode. A single number would be reassuring and wrong, so what you get is the exact byte ceiling and a progress bar.

What does bound a run, so it can never hold the panel indefinitely:

- the byte ceiling for the mode — 8192 or 32768;
- a quiet line: one second with no activity ends it;

- under flow control, 32 windows or 20 minutes of wall clock, whichever comes first. The 20 minutes is what matters on slow lines, where a full 32 kB window would take 18 minutes at 300 baud.

A tap on Capture during a run is not a stop. It is queued by the instrument and delivered when the run finishes. The app recognises presses made during the run it has just finished and discards them, saying so on the note row, rather than starting another recording you did not ask for.

2. It can time a capture

Set `Options ▶ Trigger = Trigger key`, press **Capture**, and the capture waits for you: it fires when you press the physical key, so you decide the moment the line is sampled. Verified on this hardware — armed on a deliberately idle line, the capture fired on the key press rather than timing out.

If the trigger you chose never arrives, the capture goes ahead free-running after a few seconds rather than failing, and the note row says which one was missing.

These two jobs never collide: the key times a capture only when you have asked for that in Options, and stops a run only while one is going.

How much you can capture at once

Four answers, and almost everyone only ever needs the first. **All of them are one press.**

You need	Use	Get
a look at what a device is saying	FRAME — just press Capture	~ 240 bytes , on screen
a longer burst, quickly	Mode → 8 kB, then Capture	up to 8 192 bytes , to a file
the longest single recording	Mode → 32 kB, then Capture	up to 32 768 bytes , to a file
more than that, losslessly	flow control, and a device that obeys it	as much as you like, ~20 minutes per press

FRAME is the normal case and covers almost all debugging. One press gives you a screenful, already decoded, with the baud rate and frame format worked out for you and a copy appended to the USB key automatically. You do not have to set anything up or lock anything down.

Recording to a file

Press **Mode** until the cell reads 8 kB or 32 kB. **You must lock the baud rate first** — press **Lock Rate**, or type the rate into **Options**. Then **press Capture once**, and that press does the whole job: it records, decodes every byte, writes them to the USB key, and comes back with the tail on screen. There is nothing to press in the middle.

SERIAL DECODE - 8K CAPTURE

End App

BAUD	FORMAT	IDLE	LOGIC	THRESH	SA/BIT	RATE	BYTES	ERR	FIT
19200	7E1	--	?	--	--	77 ks/s	--	--	0.00
Start bit									

8 kB decoding... 7772/8192 bytes

no stop until it ends

Capture

View

Mode

NewLog

Save


Options

While it works, the counter on the status row moves — first `recording... N % of the buffer`, then `decoding... N/8192 bytes` — with a cyan progress bar beside it, and the cell to its right says whether anything can stop it. The screenshot above is a real 8 kB run caught at 7772 of 8192 bytes decoded; the dump area is blank because nothing has been shown yet, and the note row is empty because there is nothing wrong. **There is no way to finish it early** — see the section on the TRIGGER key above for what bounds it instead.

When it finishes, the note row goes quiet. The status row's summary — bytes, errors, window count and filename — is the completion message; the yellow note row is reserved for what went wrong, so an empty one after a recording means there was nothing to say.

It stops by itself when the window is full or when the line has been quiet for a second. There is also a 20-minute backstop, so nothing can wait for ever — it only matters on very slow lines, where a full 32 kB window takes 18 minutes at 300 baud.

Nothing is dropped inside a recording. What you get is a continuous, in-order run of the line — verified with a 1024-byte non-repeating test pattern, where the decoded bytes were one unbroken slice of it.



If a **Processing reading backlog...** dialog appears, ignore it. Do not press Abort.

It is the instrument's own message, not the app's: the digitizer is producing readings faster than the meter is moving them out of the buffer. It **clears by itself** and the recording carries on normally.

Seen on the bench from 38.4 kBd upwards — at 182 kS/s, at 635 kS/s recording a 76.8 kBd line, and within a second of pressing Capture at 1 MS/s on a 115.2 kBd line. Every one of those recordings completed and decoded correctly, including a full 32 kB one.

Pressing **Abort** would stop the acquisition the dialog is describing, which is the one thing that actually loses the capture. There is nothing to do but let it run.

The photograph above is that dialog over a live 32 kB recording of a 76.8 kBd line at 332 kS/s, with the progress bar behind it at 32 % of the buffer and the row beside it reading **no stop until it ends**. It is a photograph rather than a screenshot for a reason worth knowing: **the instrument's web interface stops answering while the digitizer runs**, so the remote screen grab cannot capture this dialog at all -- eight attempts timed out. If you need to document or report something that happens during an acquisition, a camera is the only way to see it.

Which window size to pick

The window is the unit of work: one press records a window, decodes it, and files it. So the size is a straight trade, and neither choice affects what gets decoded — both are lossless.

	8 kB	32 kB
Bytes per press	8 192	32 768
Recording, at 9600 baud	~8.5 s	~34 s
Decoding afterwards	~30 s	~2 minutes
Best for	seeing something soon, and stopping soon	the most data for the fewest presses

The app starts in FRAME, and one press of **Mode** reaches 8 kB, a second 32 kB. Stop at 8 kB when you would rather see bytes soon than see all of them, or while you are still working out whether you are probing the right pin; go on to 32 kB when you want the most data for the fewest presses.

Above: an 8 kB recording has finished and **Mode** has been pressed once, so the title bar reads **32K CAPTURE** while the previous run's retained tail is still on screen. Note what the panel is saying — **Up** and **Dn** have appeared because the tail is more than one page, the status row states the next run's ceiling and that nothing will stop it, and **the note row is empty**, because a recording that worked has nothing to warn about.

More than one window: flow control

Set **Options ▶ Rear BNC = FC Out**. The rear EXT TRIG OUT then emits **one ~5 V, ~10 μs pulse each time a capture is armed** (measured on a scope: 4.92 V, 9.4–9.8 μs — the width is the meter's own default and this firmware does not let the app change it) — a credit meaning "*send now, and send no more than one window's worth*". Then one press of **Capture** runs the whole conversation: credit, record, decode, file.

credit again — round and round, with the window number on the status row, until the device stops sending. While the app is decoding, nothing is armed, no credit is issued, and a device that waits stays quiet, so the gap where bytes could be lost is closed by design. Everything lands in one file, in order.

It ends when your device goes quiet after a credit, or when it reaches one of two backstops: **32 windows** or **20 minutes**. There is no key that ends it early. Those exist so a device that never stops talking cannot leave the panel busy for ever, and the 20 minutes is usually the one you meet — decoding a window takes longer than recording it, so 20 minutes is about eight windows at 9600 baud rather than 32.

If it stops for either reason it says which, and it tells you how to carry on: **press Capture again**. A finished recording leaves you in the mode you chose — the title bar still reads `SERIAL DECODE - 32K CAPTURE` — so it is one press, and the bytes that follow go into a **new** numbered file rather than the one it was just writing.

No bytes are lost in the gap. Your device only transmits when it is credited, and no credit goes out until the next recording arms, so it is still waiting exactly where the run stopped. What you get is one conversation split across two files, in order, with nothing missing between them.

Three things your device has to get right:

- **Catch a 10 μ s pulse.** That is short: use an interrupt, an edge-triggered input or a hardware latch. A firmware polling a GPIO in a loop can miss it, and at 250000 baud 10 μ s is only 2.5 bit times. It is a pulse per window, not a level you can poll once.
- **Send no more than one window per pulse** — 8 192 or 32 768 bytes depending on the mode. The pulse carries no size information, so the limit has to be built into your firmware.
- **Go quiet when it has nothing to send.** That silence is what tells the app the transmission is over.

Honest limitation: the pulse itself is verified — 4.92 V, ~ 10 μ s, one per arm, measured on a scope — and the loop that issues them is tested. The full round trip, with a device that actually waits for the pulse, has **not** been tested here, because the signal generator on this bench cannot be gated by the meter's trigger output. The mechanism is sound by construction; you would be the first to close the loop.

How long a recording gets you

There are two different times here and it is worth keeping them apart:

- **Signal** — how much of the line ends up in the file.
- **Waiting** — how long you stand there before the buffer is full.

Below about 19200 baud they are the same. Above it they are not, because the meter writes about **100 000 readings a second** into its memory however fast the digitizer is told to sample. The samples are properly spaced and the record is complete; gathering them just takes longer than the stretch of line they represent, so you stand there longer than the signal you get.

That ceiling is a separate thing from the sample rate, and the two are easy to confuse. The **sample rate** decides how many samples land in each bit — a frame capture gets very nearly the rate it asks for, a recording's burst gets somewhat less, and that difference is what limits recording to 153600 baud (above). The **100 000 readings a second** is not about fidelity at all: it is how fast readings reach memory, and it decides only how long you wait.

Locked rate	Signal you get	Waiting for the full 32 kB
300 Bd	1092 s	the same
1200 Bd	273 s	the same

Locked rate	Signal you get	Waiting for the full 32 kB
2400 Bd	136 s	the same
4800 Bd	68 s	the same
9600 Bd	34 s	the same
19200 Bd	17 s	the same
38400 Bd	8.5 s	about 17 s
57600 Bd	5.7 s	about 17 s
115200 Bd	2.8 s	about 18 s
153600 Bd	2.1 s	about 21 s

Every row is the same 32 768 bytes — what changes is how long that takes to go past. The 8 kB window is a quarter of each figure. Then add the decoding, which does not depend on the baud rate: about two minutes for 32 kB, half a minute for 8 kB.

You do have to fill it. There is no way to stop a recording early — see the TRIGGER key section above — so choose the window with the whole wait in mind.

153600 baud is the fastest rate that can record. Above that, FRAME only.

The reason is that a recording and a frame capture do not get the same sample rate out of the same request. A frame capture gets very nearly the 1 MS/s it asks for; a recording's burst acquisition loses about half a microsecond per sample, so the same request delivers about 662 kS/s. The floor is four samples per bit, and the crossover works out at **165563 baud** — so of the standard rates, 153600 is the last one that clears it, at 4.3 samples/bit.

Above it the app **refuses before you press**, on the note row, naming the figure a recording would actually get:

192000 Bd: a recording gets 3.4 samples/bit, under the 4 floor -- use FRAME

That figure is deliberately the recording's, not the header's. The SA/BIT cell will read about 5.2 at 192000 baud, which is true of the frame capture it was measured from — the two are different acquisitions, and the note says which one it means.

FRAME keeps working to the full 250000 baud, and usefully so: at 192000 baud a frame capture holds 384 bytes.

Reading the bytes

The hex view shows two groups of eight bytes and then the same bytes as text. The offset down the left is dimmed because it is there to navigate by, not to read. A byte that could not be recovered shows as --.

ERR 0 does not mean the bytes are definitely right. It counts broken stop bits and parity failures. Noise landing inside a data bit flips it with nothing to detect — which is why the note row matters more than the error count.

If you see **began mid-byte – the first N bytes are misaligned.** : on a line that never goes quiet, the capture starts wherever it starts, and if that is halfway through a byte the first few bytes are chopped up wrongly. They are shown rather than hidden, but **do not trust their values**. Everything after the first gap in the traffic is fine. On traffic with normal gaps between messages it costs **5 to 12 bytes**; on a line that never pauses it can be far more — **up to 165 bytes measured** on back-to-back text at 250000 baud, because there is no gap to resynchronise on. A device that pauses between messages never shows it at all.

Locking the baud rate, and the amber padlock

The app locks the rate for you after a good capture, which is why later captures are quicker. It only does this when it is confident: the rate landed on a standard one, enough clean frames came through, **FIT** was 0.9 or better, and the voltage levels were steady.

When it is not confident it leaves the padlock **amber** and offers you the **Lock Rate** button. The bytes are still shown — they just are not trusted enough to carry forward. Press **Lock Rate** to pin it anyway, or capture again.

Auto-lock never locks polarity. It is worked out fresh on every capture, on purpose: guessing it wrong once and then sticking to it is how you get a confident, completely wrong decode. (**Lock Detected** on the Options screen *does* copy the detected polarity into the fields along with everything else — but there you can see what it picked before you commit it.)

If the line changes rate, the app gives the lock up

A locked rate is a promise about the wire, and when the wire stops keeping it the lock is worse than useless — it turns every frame into an error and nothing gets better by pressing Capture again. So when **all three** of these hold, the app **drops the lock by itself** and re-detects from the same samples:

- a rate is locked;
- the bit time it measures on the wire disagrees with that rate;
- more than a quarter of the frames in the middle of the capture failed.

It says so on the note row, naming the rate it discarded — **19200 baud fit nothing -- unlocked, and this capture reads 38400** — and from there auto-lock pins the new rate on the usual terms. That is the same result as typing **0** into Options by hand, which is what you had to do before.

The middle condition is what protects a rate you typed deliberately. If the rate fits the wire and the *format* is wrong, every frame fails too — but the measurement agrees with your number, so your number is kept.

Options

SERIAL DECODE OPTIONS

End App

Baud Rate

0 = auto-detect

0.00000 Bd

Data Bits

auto, or fix it

Auto (7/8)

Parity

auto or fixed

Auto

Polarity

idle level

Auto

Trigger

capture start

Start bit

Rear BNC

trig in / flow-control out

Off

Auto Detect

Lock Detected

Apply

Cancel

Setting	What it does
Baud Rate	0 means auto-detect. Type a number to lock that rate; it is rounded to a whole baud.
Data Bits	Leave on Auto (7/8) . Only use Auto (any) if you know the device sends 9-bit words — it can make a damaged capture look tidier while being wrong.
Parity	Auto , or pin None / Even / Odd .
Polarity	Auto , Idle high (normal CMOS/TTL) or Idle low (inverted). Auto is reliable unless the line is almost never idle.

Each setting is independent — pinning one leaves the rest on auto.

Apply commits your changes and takes a capture. **Cancel** puts everything back to how it was when you opened the screen. **Auto Detect** forgets everything locked, takes a fresh capture, and locks whatever that finds — it is the quick way back when you have locked onto the wrong thing. **Lock Detected** copies the last capture's settings into the fields so you can look at them before committing.

There is no stop-bit setting. A second stop bit looks exactly like a bit of idle line, which the decoder already copes with, so a 2-stop device decodes correctly and is reported as 1 stop.

Advanced

Two settings you can ignore unless you need them:

- Trigger** — what starts the acquisition. **Start bit** (the default) waits for the line to move. **Free run** grabs immediately. **Trigger key** waits for the front-panel TRIGGER key, as described under **The TRIGGER key** above. **If the thing you asked for never happens, the capture goes ahead anyway** after a few seconds and the note row names the one that did not arrive — **edge trigger**

unavailable; captured free-running, or front for the TRIGGER key. Those bytes are real but they are not lined up with the event you wanted.

- **Rear BNC** — lets an external trigger in on the rear connector, and FC Out sends the flow-control credit pulse out of EXT TRIG OUT. FC Out is what turns a recording into an unlimited one; see **More than one window: flow control** above.

This setting affects FRAME captures. **A recording ignores it** and starts as soon as you press Capture — for a recording, the thing that decides when the device talks is the flow-control credit, not a trigger.

Files on the USB key

Two different files, and they are not the same thing:

File	Written	Holds
/usb1/bytesNNN.txt	automatically, on every capture	the running log — every capture's bytes as text and hex. NewLog starts a new number.
/usb1/serial_NNN.txt	only when you press Save	a full report of what is on screen: settings, all the notes, the whole dump.

A recording decodes into the log file, so the file holds the whole run while the panel shows only the tail of it. The note row says so when they differ.

Recording needs a USB key — the file is the point of it. Normal FRAME captures work without one; they just cannot log or Save.

When something looks wrong

What you see	What it means
line is idle (no transitions)	Nothing is crossing the threshold. Check your connection, and check the device is actually talking while you capture.
no clear logic levels	Too much noise, or the swing is too small. 1.6 V works; 60 mV is refused outright.
Wrong rate reported	The traffic may be genuinely ambiguous (below). Lock the rate if you know it.
Bytes wrong but ERR 0, with the rate locked	The device is not running at the rate you locked. The note row warns when the line disagrees by more than 4 %.
A baud rate that is not a round number	It is reporting what it measured, because the measurement was not within 2 % of a standard rate.

Some traffic is genuinely ambiguous, and no decoder can fix it. Eight 0x00 bytes at 9600 baud 7N1 make exactly the same waveform as four 0x08 bytes at 4800 baud 8N1. The app tells you a rival rate fits rather than pretending to know. Irregular gaps between bytes remove the ambiguity completely.

Two cases where automatic detection is known to fail are described under **Known failures of automatic rate detection**, below. Both are fixed by typing the rate in.

What it can cope with

Measured at 8 samples per bit. Past every one of these limits the app **reports** the problem rather than quietly getting it wrong, except where noted.

Condition	Works up to
Wrong clock, rate locked	$\pm 4\%$. Beyond that a fast device corrupts bytes with nothing to show it
Wrong clock, auto-detect	$\pm 12\%$ and past — it measures instead of assuming
Jitter	$\pm 15\%$ of a bit; past $\pm 20\%$ it can fail quietly
Noise	about 40% of the logic swing — roughly 1.3 V on a 3.3 V line
Slow edges, long cables, weak pull-ups	filtering up to 0.7 bit times
Logic swing	1.6 V up to 5 V TTL
DC offset, logic low at or above 0 V	anything that keeps the line inside ± 10 V
DC offset, logic low below 0 V	works, but see the note below — one band is not reliable

Rates from 300 to 250000 baud decode byte-exact on the bench, and so do odd ones — 900, 1500, 3600, 7200 and arbitrary values like 1379, 8123 and 104857. Two exceptions are known and both are **flagged rather than silent**: long gapless payloads at 2400 baud produce one to three damaged frames in a few hundred, and one 250000-baud run matched 497 of 499 bytes with the odd byte marked. A rate within 2 % of a standard one is reported as the standard one, because real devices use standard rates.

Tolerance follows samples per bit, not baud rate. A slow line is not automatically the safe one — read `SA/BIT`.

If you are writing the firmware at the other end, two things help: keep your edges within $\pm 15\%$ of a bit time (that is $\pm 15.6\ \mu\text{s}$ at 9600 baud but only $\pm 1.3\ \mu\text{s}$ at 115200, so if an interrupt can delay your transmit loop, send slower), and avoid padding with `0xFF`, which gives the rate detector only one narrow pulse per frame to measure. `0x00` padding is fine.

Known failures of automatic rate detection

Read this in proportion. Most serial decoders require you to tell them the baud rate. This one measures it, and gets it right on every standard rate and on arbitrary ones like 1379, 8123 and 104857. Two cases are known where the measurement goes wrong, together fewer than **one point in 200**, and in both of them the remedy is to supply the rate — which is what other tools ask for as a matter of course.

Both are fixed by typing the rate into `Options > Baud Rate` — verified, every frame decoded with zero errors once the rate was given. The numbers are below so you can judge the risk yourself.

Automatic detection has no clock to refer to. It works in two steps: it measures the gaps between transitions and finds the longest bit time that divides them all, then weighs that answer against the standard rates, because a real device almost always uses one.

Gaps of many different lengths pin the measurement down well; gaps of only a few lengths do not. But where detection fails it is usually not the measurement that is wrong — it is the second step. A different reading at a simple multiple of the measured rate can land on a standard value while the true rate does not, and that coincidence is allowed to outweigh a correct measurement. **This is why the failures below happen only on non-standard rates:** when the true rate is itself a standard one, nothing can outbid it.

1. A short pattern repeated over and over

What triggers it. Data built from `0x00`, `0xFF`, `0x55`, `0xAA`, walking bits, or one fixed string sent on a loop — link-training patterns and idle fill, rather than real messages. Such data yields only a handful of distinct gap lengths, and a simple fraction of the true bit time then explains those gaps just as well as the truth: half the bit time reads the same square wave as twice as many bits, with no framing error anywhere to give it away. Both readings describe the waveform, so nothing in the signal can rank them — and where one of them lands on a standard rate and the true rate does not, the standard one wins.

The measurement is not the problem. Traced on a failing case, the bit time was recovered as 8.3990 samples against a true 8.3987, every one of the smallest gaps sat at the true bit time, and the decode at that rate framed every byte without a single error. It was then overruled by its own half, because 3572 is not a standard rate and 7200 is.

What you see. The reported rate is a small whole-number multiple of the true one — two to five times — and always one of the standard rates. Expected 3572 and the panel says 7200, or expected 124000 and it says 250000. Usually the byte area also fills with error markers, because a wrong bit time shreds most frames; in the measured cases 83 to 119 frames of a few hundred were flagged. Occasionally the frames survive and the bytes are simply wrong, which is the case worth knowing about.

There is no warning about the rate itself. Where two rates fit a waveform the panel normally names the runner-up under the reading, but not in this case — verified, the note is empty here. So the error markers are the whole of the signal, and where the frames happen to survive there is none. If a repeating test pattern reports a standard rate you did not configure, distrust it and lock the rate.

How rare. 24 of 6 714 capture-and-decode points across four full bench sweeps — **0.36 %**, **about one in 280**. Every one was a repeating pattern at a non-standard rate. **No ordinary payload at a standard rate has ever shown it** — the fox, the 1 kB text payload and twelve random payloads are clean across every rate from 300 to 250 000 baud.

Split by rate, it is not spread evenly at all. Over a separate 12 341-point sweep, seven laps across 41 waveforms and 43 rates:

rate	points	affected	
a standard baud rate	6 314	0	0.000 %
a non-standard rate	6 027	24	0.398 %

Not one failure at a standard rate **in that offline sweep**, and there is a structural reason: a bit time that lands on a standard rate and decodes without error is not overruled, because the check that would allow it requires the reading to be visibly failing first. The worst case there is the opposite — a rate that is an exact half of a standard one without being standard itself: **125 000 reported as 250 000 accounts for 7 of the 24**.

On the instrument itself, however, standard rates are not immune, and one case is worth stating plainly. A waveform carrying a **LIN-style break field** — 13 or more bits held low, which is not valid 8N1 at all — was measured at **three times** its true rate at 600, 19 200, 31 250, 38 400, 57 600 and 76 800 baud, in every lap of a 14-hour soak, and **with no framing errors to warn of it**. Expected 19 200, reported 57 600, and the bytes were confidently wrong. This decoder does not claim to read LIN, and a break field is outside UART by construction — but if you are probing a bus that begins frames with a long dominant field, set the rate by hand and do not trust the measured one.

No **ordinary UART payload** at a standard rate has shown this: the fox, the 1 kB text payload and twelve random payloads are clean at every rate from 300 to 250 000 baud, on the bench and offline.

One qualification, which is not this failure. Half of a standard rate is often *also* standard — 4800 is half of 9600 — and for perfectly regular traffic the two readings are the same waveform, as **Some traffic is genuinely ambiguous** above describes. That is not a misdetection and the panel names the rival rate when it happens; the failure in this section is the one where it does not.

The bench deliberately spends **half** its points on non-standard rates, which is far more than ordinary use, so the figure above overstates the exposure of a typical setup by roughly five times.

Besides locking the rate: capture real traffic rather than a test pattern, or let the bytes have irregular gaps — either removes the ambiguity.

2. Logic low sitting well below ground

What triggers it. A line whose low level is *negative* with respect to the probe's reference, with the whole signal still inside ± 10 V.

What you see. The rate comes back wrong, in the same multiple-of-the-truth way as above.

How rare. Most of the negative range is fine: measured clean at **-0.1, -0.25, -0.5, -1.0 and -3.0 V** on a 3.3 V swing. Around **-2 V** some payloads fail. It is not a threshold — **-3 V works where -2 V does not** — so it is a band rather than a limit. Across a full seeded sweep with amplitude and offset both varied, **12 of 1763** points were affected (**0.68 %**), and varying the amplitude alone over 0.5× to 1.6× caused **no failures at all**. Signals sitting entirely at or above 0 V are unaffected at any offset inside ± 10 V.

Besides locking the rate: reference the probe to the line's own low level rather than to earth.

Measured figures, the full tolerance envelope, press-by-press timings and the reasoning behind the design are in [REFERENCE.md](#).